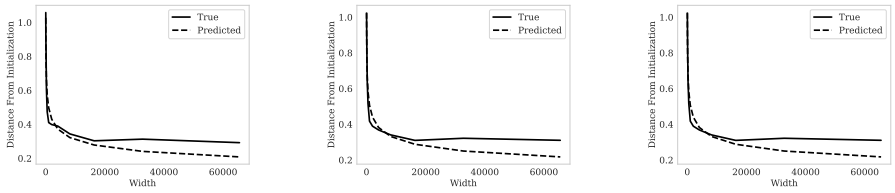

(a) .1 Dropout Probability (b) .2 Dropout Probability (c) .3 Dropout Probability

Figure 1: Distance from Initialization for with Dropout for different Dropout probabilities.

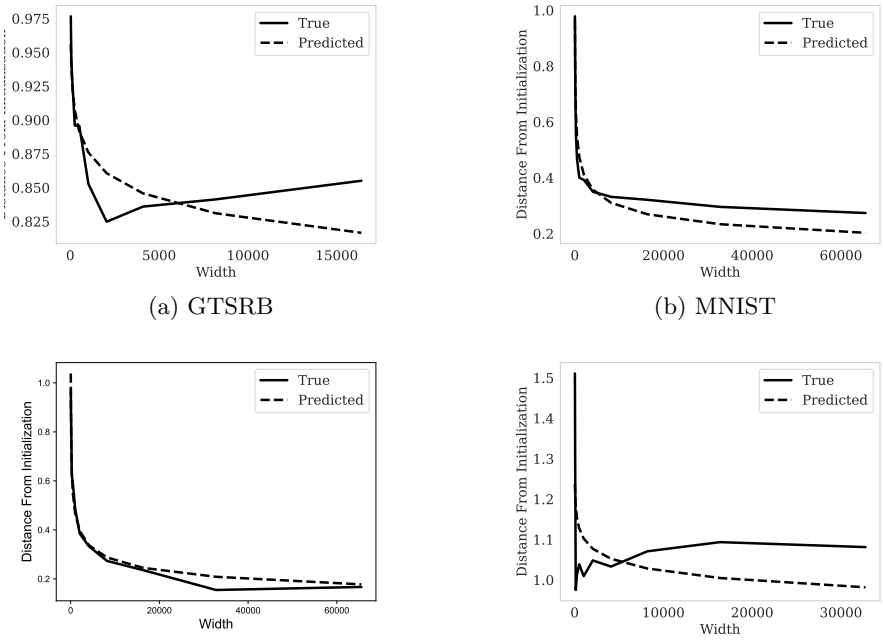

(a) GTSRB (b) MNIST

(c) Fashion MNIST (d) SVHN

Figure 2: We plot the distance from initialization for training at different widths on different datasets. In the "Predicted" line, we plot the approximation from Assumption 4.3.

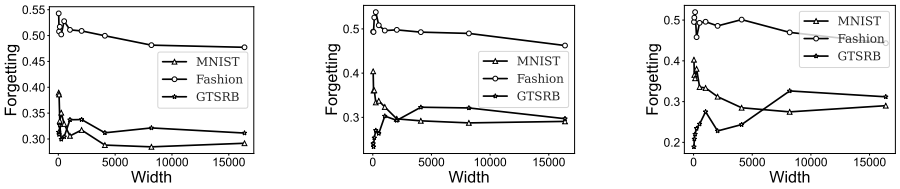

(a) 4 Hidden Layers (b) 5 Hidden Layers (c) 6 Hidden Layers

Figure 3: We plot forgetting as width of the network is increased for different number of hidden layers.

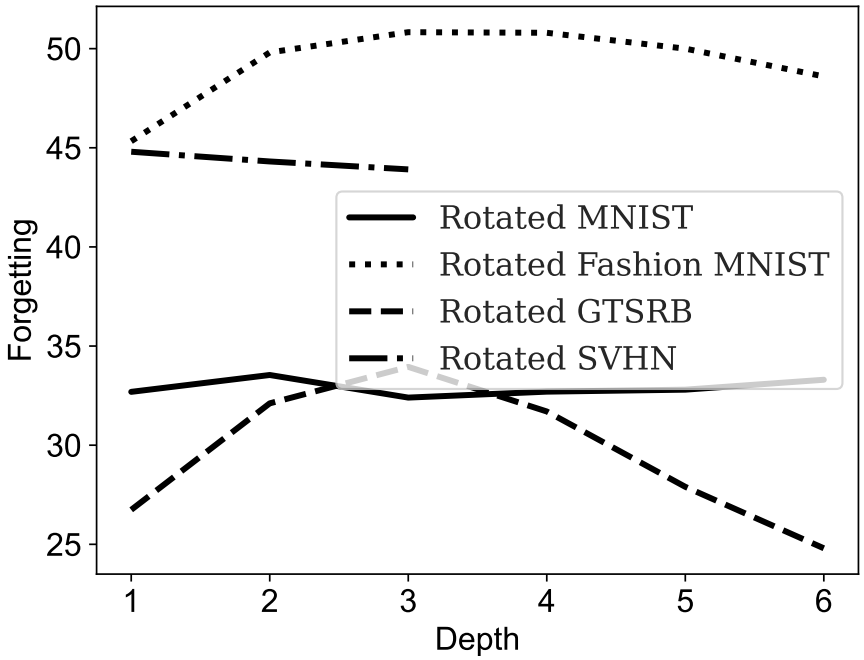

Figure 4: Average Forgetting as Depth is increased. We see that as depth is increased, forgetting increases. However, as depth is increased further, the accuracy goes down due to vanishing gradients. This artificially causes the forgetting to decrease.